

Question 1

Display the data types of each column using the function `dtypes`, then take a screenshot and submit it, include your code in the image.

```
In [7]: df.dtypes

Out[7]: Unnamed: 0      int64
id          int64
date        object
price       float64
bedrooms    float64
bathrooms   float64
sqft_living  int64
sqft_lot     int64
floors       float64
waterfront   int64
view         int64
condition    int64
grade        int64
sqft_above   int64
sqft_basement int64
yr_built     int64
yr_renovated  int64
zipcode      int64
lat          float64
long         float64
sqft_living15 int64
sqft_lot15   int64
dtype: object
```

We use the method `describe` to obtain a statistical summary of the dataframe.

```
In [7]: df=pd.read_csv(file_name)

df.drop(["id", "Unnamed: 0"], axis=1, inplace = True)

df.describe()
```

```
Out[7]:
```

	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	condition	
count	2.161300e+04	21600.000000	21603.000000	21613.000000	2.161300e+04	21613.000000	21613.000000	21613.000000	21613.000000	21613.000000
mean	5.400881e+05	3.372870	2.115736	2079.899736	1.510697e+04	1.494309	0.007542	0.234303	3.409430	7.656870
std	3.671272e+05	0.926657	0.768996	918.440897	4.142051e+04	0.539989	0.086517	0.766318	0.650743	1.175450
min	7.500000e+04	1.000000	0.500000	290.000000	5.200000e+02	1.000000	0.000000	0.000000	1.000000	1.000000
25%	3.219500e+05	3.000000	1.750000	1427.000000	5.040000e+03	1.000000	0.000000	0.000000	3.000000	7.000000
50%	4.500000e+05	3.000000	2.250000	1910.000000	7.618000e+03	1.500000	0.000000	0.000000	3.000000	7.000000
75%	6.450000e+05	4.000000	2.500000	2550.000000	1.068800e+04	2.000000	0.000000	0.000000	4.000000	8.000000
max	7.700000e+06	33.000000	8.000000	13540.000000	1.651359e+06	3.500000	1.000000	4.000000	5.000000	13.000000

we can see we have missing values for the columns `bedrooms` and `bathrooms`

```
In [8]: print("number of NaN values for the column bedrooms :", df['bedrooms'].isnull().sum())
print("number of NaN values for the column bathrooms :", df['bathrooms'].isnull().sum())

number of NaN values for the column bedrooms : 13
number of NaN values for the column bathrooms : 10
```

Question 3

Use the method `value_counts` to count the number of houses with unique floor values, use the method `.to_frame()` to convert it to a dataframe.

```
In [27]: df['floors'].value_counts().to_frame()
```

```
Out[27]:
```


	floors
1.0	10680
2.0	8241
1.5	1910
3.0	613
2.5	161
3.5	8

Question 4

Use the function `boxplot` in the seaborn library to determine whether houses with a waterfront view or without a waterfront view have more price outliers.

```
In [29]: sns.boxplot(x='waterfront', y='price', data=df)

Out[29]: <AxesSubplot:xlabel='waterfront', ylabel='price'>
```




Question 5

Use the function `regplot` in the seaborn library to determine if the feature `sqft_above` is negatively or positively correlated with price.

```
[30]: sns.regplot(x="sqft_above", y="price", data=df, ci = None)

Out[30]: <AxesSubplot:xlabel='sqft_above', ylabel='price'>
```



We can use the Pandas method `corr()` to find the feature other than price that is most correlated with price.

```
[ ]: df.corr()['price'].sort_values()
```

Question 6

Fit a linear regression model to predict the `'price'` using the feature `'sqft_living'` then calculate the R^2 . Take a screenshot of your code and the value of the R^2 .

```
[34]: Z = df[['sqft_living']]
      Y = df['price']
      lm1 = LinearRegression()
      lm1.fit(Z,Y)
      lm1.score(Z,Y)

Out[34]: 0.4928532179037931
```

Question 7

Fit a linear regression model to predict the `'price'` using the list of features:

```
[5]: features = ["floors", "waterfront","lat" ,"bedrooms" ,"sqft_basement" ,"view" ,"bathrooms","sqft_living15","sqft_above","grade","sqft_living"]
```

Then calculate the R^2 . Take a screenshot of your code.

```
[8]: XX = df[features]
      YY = df['price']
      lm2 = LinearRegression()
      lm2.fit(XX,YY)
      lm2.score(XX,YY)

[38]: 0.6576569675583581
```

Question 8

Use the list to create a pipeline object to predict the 'price', fit the object using the features in the list `features` , and calculate the R^2 .

```
In [40]: pipe=Pipeline(Input)
         pipe
         pipe.fit(XX,Y)
         pipe.score(XX,Y)

Out[40]: 0.7513417707683823
```

Question 9

Create and fit a Ridge regression object using the training data, set the regularization parameter to 0.1, and calculate the R^2 using the test data.

```
In [44]: from sklearn.linear_model import Ridge
```

```
In [45]: RM = Ridge(alpha=0.1)
RM.fit(x_train, y_train)
RM.score(x_test, y_test)
```

```
Out[45]: 0.6478759163939113
```

Question 10

Perform a second order polynomial transform on both the training data and testing data. Create and fit a Ridge regression object using the training data, set the regularisation parameter to 0.1, and calculate the R^2 utilising the test data provided. Take a screenshot of your code and the R^2 .

```
[46]: pr=PolynomialFeatures(degree=2)
x_train_pr=pr.fit_transform(x_train[features])
x_test_pr=pr.fit_transform(x_test[features])
```

```
RM2 = Ridge(alpha=0.1)
RM2.fit(x_train_pr, y_train)
RM2.score(x_test_pr, y_test)
```

```
Out[46]: 0.7002744273468813
```